

Helios: Your Private Interview Coach

Setup Guide

Get ready for offline practice in 5 simple steps.

Based on Helios v0.1.0 · Local-First Architecture

Section 1 — Ready Status Checklist

All five core components are installed and ready to go. Here's your at-a-glance summary:

	Language Model Ollama · qwen3:8b	 Ready
	Transcription Whisper · faster-whisper	 Ready
	Speaking Voice Piper · en_US-lessac	 Ready
	Voice Detection Silero VAD	 Ready
	Speaker Separation Pyannote · (optional)	 Ready

Section 2 — Step-by-Step Configuration

■ Step 1 — Start the Engine

Double-click **Start-Helios.cmd** in your Helios folder. A terminal window will open and begin loading all components. Wait until you see:

■ Application startup complete

Leave this window open — it keeps Helios running. You can minimise it.

■ Step 2 — Connect the Brain (AI Model)

Open Helios in your browser, go to **Settings** → **Your AI connection**, and configure the three highlighted fields exactly as shown:

■ Your AI connection

Offline mode

ON

Provider

Ollama – local

Model

qwen3:8b

Cloud fallback

— disabled —

All inference runs locally. No data leaves your machine.

■ Pro Tip — Download the Model First

Before starting Helios, open a terminal and run:

```
ollama pull qwen3:8b
```

This downloads the AI model (~5 GB). It only needs to happen once.

■ Step 3 — Give it a Voice & Ears

In **Settings** → **Transcription** and **Settings** → **Speaking Voice**, click the **"Set up model"** button for each. The models download automatically — no extra steps needed.

■ Transcription	faster-whisper (base.en)	Click Set up model
■ Speaking Voice	en_US-lessac-medium	Click Set up model

■ Step 4 — Advanced Features (Optional)

Speaker Separation (Pyannote) lets Helios label who is speaking — useful if you're practising with a partner. It is optional and disabled by default.

■ Before enabling Speaker Separation:

1. Create a free account at huggingface.co.
 2. Accept the Pyannote model licence conditions on the model page.
 3. Generate an Access Token and paste it into Helios Settings.
- You can skip this entirely and everything else works fine.

Section 3 — Tuning Your Experience

■ Rhythm of your conversation — Pause detection slider

This slider controls how long Helios waits after you stop speaking before it considers your turn finished. A shorter pause makes the AI respond quicker; a longer pause gives you more time to think mid-sentence.

1.0 s	Very quick — good for rapid Q&A
1.8 s ← Recommended	Natural pacing for interview responses
3.0 s	Relaxed — ideal for long detailed answers

■ **Headphone Tip — Prevent Echo**

Using speakers can cause Helios to hear itself and create an echo loop. To prevent this, either use headphones (recommended) or enable **Ignore microphone while speaking** in Settings → Microphone. This mutes the mic whenever the AI is talking.

Section 4 — Your Privacy Promise

■ **Your Data Stays Yours — Always.**

Helios is built on a local-first architecture. Here is exactly what that means for you:

- **No cloud uploads.** Your voice, transcripts, and sessions never leave your device.
- **No telemetry.** Helios does not phone home or collect usage statistics.
- **Local SQLite storage.** All session history is saved in a database on your own computer.
- **Secure keyring.** If you add an API key (for cloud fallback), it is stored in your OS keychain — not in a plain-text file.
- **100% offline inference.** With Ollama + Offline Mode on, there is no cloud fallback. Period.